

Falsification-First Decipherment: A Decontaminated Inference Framework for Testing Undeciphered-Script Claims, with Linear A as a Worked Null

Anonymous TACL submission

Abstract

Decipherment research has a false-positive problem and no shared instrument: a small corpus, an unknown language, and many candidate readings let an analyst “translate” almost anything. The flagship survey lists fifteen challenges of script and data; none addresses the orthogonal inference-control layer — overfitting, multiple testing, contamination — a candidate sixteenth challenge. We build a falsification-first framework for it: a fabricated-language negative-control family, multiple-comparison deflation, a literature index separating published from not-indexed signs, and pre-registration with a mechanically-graded gate. This is a methods paper; we assert no decipherment. The machinery calibrates: on a tested null the gate’s false-graduation rate is 0.6% (3/500); the identical morphology test finding no power on Linear A recovers Mycenaean (Linear B) inflection (9/16 affixes above the bigram floor), supporting the interpretation that Linear A’s short forms contribute materially to its lack of power. Across probes it returns differentiated, calibrated outcomes — null, no-power, supported-structural, circular, and contamination-sensitive. Results are consistent with identifiability, rather than corpus size, being the more fundamental present constraint. The June-2026 *301 “cracked” claim is a qualitative illustration of the failure modes the gate is designed to detect.

1 Introduction

Decipherment has a false-positive problem, and the field has no agreed instrument for measuring it. The history of undeciphered-script research is, in large part, a history of uncontrolled model selection: a tiny corpus, an unknown language, and a large space of candidate readings let a determined

analyst “translate” almost anything, with internal consistency on the cherry-picked set standing in for confirmation. The mechanism is not any single bad match but uncorrected multiplicity: with enough candidate roots and a small enough corpus, a pile of individually-plausible matches is the *expected* outcome under pure chance, so “each individual comparison looks suggestive [while] the conclusion is nonsense”. A standard demonstration of spurious comparison: one can “prove” *English is a Semitic language* by matching a handful of English words to Proto-Semitic roots while quietly ignoring the far larger number that do not match. The same move, applied to a real undeciphered script, has produced a long line of decipherment claims that no one can either confirm or refute.

That the discipline lacks a shared standard here is not our editorializing but the field’s own assessment. The flagship computational-decipherment survey — Braović et al.’s (2024) systematic review of Bronze Age Aegean and Cypriot scripts — enumerates fifteen “major challenges,” from unknown writing system and unknown language to small dataset, incomplete vocabulary, allographs, and unavailable parallel data (Braović et al., 2024). Those fifteen are primarily properties of the *object and its data* — the script, the language, the surviving corpus, and the resource/comparison-data conditions (the list includes resource and hardware items too). What governs whether a *claimed reading* is believable is a different kind of thing: a property of the inference methodology and its epistemics — overfitting, multiple testing, false-positive control, and contamination. We therefore frame this axis, once and modestly, as a candidate **sixteenth challenge** — a deliberate rhetorical foil to the survey’s fifteen, not a claim to a new taxonomic slot or to sole discovery. It is *under-instrumented* rather

than unrecognised: Packard’s (1974) nine deliberately fabricated decipherments already exercised exactly this false-positive logic as a control, and the selective-inference machinery we import (Bailey and López de Prado, 2014) is mature elsewhere — our contribution is to *consolidate and operationalize* that axis for decipherment, not to invent it. It is the methodological axis the object-level taxonomy was never meant to cover, and it is distinct from the survey’s own acknowledged gap — that on scoring results “in many cases there is nothing to compare them against, and their evaluation remains an open problem” (Braović et al., 2024, Sec. 6.3, p. 746). That is an *evaluation-target* problem (what to measure against); the false-positive / multiplicity problem (whether a fitted reading is being mistaken for a real one) is separable from it, and it is the latter this paper instruments.

Contribution. We build the machinery for that axis: a falsification-first, decontaminated inference framework for testing undeciphered-script decipherment claims (Braović et al., 2024). The harness applies fail-closed validation to every positive and enforces four concrete defences: a fabricated-language canary that supplies a **generator-conditional negative-control distribution** (a false-positive floor conditional on the fabricated-language generator — here a first-order sign-bigram Markov model; a broader control ensemble spanning unigram, root-template, and real-language generators is disclosed as future work, not built here); multiple-comparison deflation for the number of signs, roots, and families a search was free to try; a literature index that partitions published (L_{known}) from literature-unseen ($L_{\text{not_indexed}}$) signs so that reproducing a known value cannot masquerade as discovery; and pre-registration with a held-out gate whose verdict is computed mechanically, never by the model that proposed the reading. The evaluation gap gets a concrete, if power-limited, answer: the recurring libation *formula* — roughly 50 inscribed objects spread across peak sanctuaries and caves at more than five findspots — supports a natural leave-one-site-out cross-validation graded from held-out data. We are explicit that with so few formula sites this cross-validation has low statistical power, and we report it as such rather than as a decisive test (Braović et al., 2024).

This is a methods paper, and Linear A is its

worked null. We do not claim a decipherment. The framework is designed to withhold validation when a proposed reading does not clear pre-registered negative controls, multiplicity adjustment, contamination checks, and held-out evaluation; when nothing survives, the calibrated null — what the corpus can and cannot support — is itself the deliverable. The working conclusion on current data is that the results are **consistent with identifiability** being the more fundamental present constraint — not a demonstrated corpus-size threshold: there is no known related language to map onto, the multiple-testing burden of candidate readings goes uncorrected, and the sign inventory is mixed and partly lossy. We do not claim additional data would be irrelevant — a null can arise from identifiability, low power, or corpus size alike. A lossy, non-injective sign-to-sound channel need not admit a decipherment at any corpus size, so we make no claim that “there is enough text to pin a reading”; the rigorous demonstration of *that* constraint is the result the field can use. The June-2026 AI-assisted “Linear A cracked” claim — reading the *301-anchored libation verb as an extinct Semitic form — serves throughout as a qualitative illustration of the failure modes the gate is designed to detect, analysed from its public, archived presentation as a claim, not a person, and neither reproducible nor held-out validated. The artifact is open and archived as an anonymized artifact (code and pre-registrations released on acceptance).

Roadmap. Section 2 situates logos against the Braović et al. taxonomy and the computational-decipherment lineage, including the cognate-ID baseline. Section 3 specifies the harness — canary, deflation, literature index, and the pre-registration/held-out gate. Section 4 develops the information-floor diagnostic and states its identifiability limits. Section 5 works the June-2026 *301 claim as a case study of the failure modes. Sections 6–8 report the pre-registered nulls: morphology (§6); metrology, phonology, and cross-script signals (§7); and contamination and sufficiency probes (§8). Section 9 draws the discussion, limitations, and conclusion.

2 Related work

2.1 Position relative to the survey taxonomy

The reference frame for computational work on the Bronze Age Aegean scripts is Braović et al.

(2024, *Computational Linguistics* 50(2):725–779), the field’s first survey focused specifically on Aegean and Cypriot scripts; its only comparable predecessor is Ferrara and Tamburini (2022), and the broader machine-learning-for-ancient-languages landscape is surveyed by Sommerschild et al. (2023, *Computational Linguistics* 49(3)). Braović et al.’s Section 4 enumerates fifteen “major challenges” — from unknown writing system, reading direction and language, through small dataset, allographs and incomplete vocabulary, to unavailable parallel data and hardware — **none** of which concerns overfitting, multiple testing, false-positive control, contamination, or distinguishing a *fitted* decipherment from a *real* one. The same gap runs through Ferrara and Tamburini’s (2022, *Lingue e Linguaggio* XXI.2:239–259) §3.5, which reviews the cognate-matching lineage logos builds on — Snyder et al. (2010) on Ugaritic↔Hebrew (the 29/30-sign benchmark our canary calibrates against), Berg-Kirkpatrick and Klein (2011), and Luo et al. (2019, 2021) — and frames Linear A, via Gelb and Whiting’s typology (Gelb and Whiting, 1975), as a Type III script (unknown script *and* language) whose “necessary validation” is “altogether impossible,” yet nowhere treats how a match is to be *deflated* for the multiplicity it was selected from. Two surveys, the same unfilled axis. We do not score this as an oversight: their challenges are properties of the object and its data, whereas the axis logos foregrounds is a property of inference *methodology* — named-but-unfilled rather than absent. logos supplies the machinery to police it: preregistration, a decontaminated derivation/held-out partition, multiplicity deflation, and a graduation gate computed mechanically rather than by the model that proposed the reading.

Braović et al. do gesture at a related but separable difficulty: evaluation. Because computational analyses of an undeciphered script have “nothing to compare them against,” they conclude (Sec. 6.3) that “their evaluation remains an open problem.” That is a distinct problem from overfitting and multiplicity control, and logos addresses it separately. Its held-out design is a concrete response — the libation formula recurs across a small set of peak-sanctuary and cave findspots, permitting a leave-one-site-out cross-validation adjudicated mechanically by `verdict.py` — though we flag at the outset that so few formula findspots give this

particular test low statistical power (quantified in Section 8). We locate our components in the survey’s own method taxonomy (Sec. 7.1) and, where prior art exists, defer to it rather than claim novelty.

2.2 The current philological synthesis

Salgarella (2025), *Writing in Bronze Age Crete: “Minoan” Linear A* (Cambridge Elements), by the author of the SigLA database, is the most current authoritative one-author overview. It concludes that Minoan is a language **isolate** — it “does not belong in any known language family” (the families used for comparison being Indo-European, Semitic and Afro-Asiatic) — typologically agglutinative with a provisional Verb–Subject–Object order, and it records that “no bilingual text has so far been found” in a corpus of roughly 2,500 inscriptions. Crucially, Salgarella ranks the etymological/comparative method below combinatorial-contextual analysis and names **digital/statistical analysis as “the next desideratum in the field.”** logos positions itself as exactly that desideratum — but under discipline, treating the isolate verdict as the binding *identifiability* constraint rather than a target to be argued away.

2.3 Prior computational attempts and the search trap

The seminal automatic-decipherment result is Luo et al. (2019), which recovers a character mapping from a lost script to a *known related* language via neural minimum-cost flow. It auto-deciphered Ugaritic→Hebrew (+5.5% over prior state of the art) and translated 67.3% of Linear B→Greek cognates, but did not address Linear A and would fail there by construction: with no known cognate plaintext, the objective has nothing to optimise toward. This cognate-language precondition is the empirical form of our identifiability argument. The cognate-ID baseline logos actually runs is *not* a reimplement of Luo’s neural flow: it is a published matcher, Tamburini’s (2025) CSA_OptMatcher (coupled simulated annealing), which reports the same Luo et al. (2019) accuracy metric. We adopt it as the non-neural baseline against which any Linear A cognate claim must clear.

Against that disciplined baseline stand the in-domain instances of the search trap: the dictionary-fishing programmes of Eu et al. (2019) and Loh and Perono Cacciafoco (2020), the lat-

ter explicitly a “brute force attack” on dictionaries. These are the concrete Linear-A cases the multiplicity correction is built to deflate. On the visual side, cross-script comparison is *not* our novelty: Daggumati and Revesz (2019, 2023) already applied CNN+SVM cross-script similarity, and Corazza et al. (2022) (“Sign2Vec”) clustered signs unsupervised; our defensible contribution is only the controlled image-versus-sequence *asymmetry* under one harness.

2.4 The discipline lineage

Two older strands supply the honesty machinery. Packard (1974), *Minoan Linear A*, constructed nine deliberately fictitious decipherments as a null control — and the genuine Ventris-value transfer beat them only about 2:1 — anticipating our permutation-null canary. The statistical lineage is selective inference and multiple-comparison correction (Bailey and López de Prado, 2014), the correction for the number of hypotheses $\times$ signs $\times$ roots $\times$ families tried that turns a raw match-rate into a defensible one; its transfer from the finance/backtest setting to script decipherment is specified in Section 3. logos’s empty cell is the intersection none of these occupy: a falsification-first, decontaminated inference framework that couples a published cognate-ID matcher (Tamburini, 2025, reporting the Luo 2019 metric) with preregistration, multiplicity deflation and mechanical held-out verdicts, released openly (as an anonymized artifact; code and pre-registrations released on acceptance), with Linear A worked as the honest null rather than a claimed crack.

3 Methods: the discipline harness

The harness is a comparison-layer pipeline that a reader can reimplement. A candidate reading enters as a *proposal* and is subjected, in fixed order, to five gates: pre-registration, deflation, decontamination, held-out graduation, and a cross-cutting rule that every gate is scored from persisted artifacts rather than from any model’s narration of its own success. Model-derived signals — LLM theses, learned embeddings, cognate-matcher assignments — enter capped at confidence ≤ 0.75 , structurally unable to *decide* a verdict; only mechanically-verified held-out reads rank higher. The design targets the false-positive / multiplicity / contamination axis that Braović et al. (2024)’s fifteen-challenge taxonomy names

but leaves unfilled (§1–2).

3.1 Pre-registration

Before any model is built or run, the predictions are frozen. The morphology pre-registration fixes, for each hypothesis, the claimed affix, its source, the falsifiable test, and the acceptance threshold — with the file’s git commit — externally anchored by the immutable artifact deposit ([artifact DOI withheld for review]) — serving as the pre-specification timestamp, so the model cannot be tuned to them after the fact. Each pre-registered affix (drawn from Thomas (2020) and Duhoux/Valério) must recur on ≥ 2 distinct stems across independent inscriptions, above a within-form permutation null, and survive multiplicity correction at a target of $p < 0.01$ corrected. Outcomes are enumerated in advance — CONFIRM, EXTEND, REFUTE/NULL, or NO POWER — so that a negative is a publishable result, not a failed experiment. Any change is a new dated pre-registration, never an edit.

3.2 Deflation

Internal fit on the derivation set is treated as no evidence. Every match rate is deflated for everything tried: $n_{\text{hypotheses}} \times n_{\text{signs}} \times n_{\text{roots}}$. The nulls are pipeline-level — a frequency-banded permutation null (after Packard, 1974) and an Euler- γ order statistic ($\text{expected_max_order_stat}(\mu_0, \sigma_0, N)$) that gives the bar a normal-approximation to the maximum over N candidates tried that the best should clear by chance. Multiplicity is instrumented, not hand-passed: a deduplicating SearchLog reports the exact distinct-candidate count n_{trials} (the deduplicated count of candidate readings attempted) a scanner produces (invariant 12). The strong structural statistic, S_{morph} , earns credit only through a productivity gate — an affix must attach to ≥ 2 distinct residual stems, so a single formula word copied across inscriptions scores zero. (Its within-form-shuffle false-positive calibration and its saturated-power point are reported in Appendix A.2.)

3.3 Decontamination

A large open-weight model has read Gordon, Best, and Di Mino, so a proposed correspondence that merely reproduces a *published* value is shared-source contamination, not independent evidence.

Three mechanisms address this.

The **literature index** records published Linear A sign-value and correspondence claims; any proposal matching the index is tagged `literature_match` and quarantined — it may be tested but can never count as discovery. The index partitions the sign inventory into `L_known` (referenced by ≥ 1 published proposal) and `L_not_indexed` (untouched). Discovery claims may rest only on `L_not_indexed` held-out success, since regurgitation can only return what is in the literature. (The conservative seed, the adversarial population, and the current census count are given in Appendix C.5.)

The **L_fake canary** is a phonotactically-plausible invented lexicon, calibrated to a real candidate language’s phoneme frequencies, root-template structure, and lexicon size, but never published and never in any training corpus. Running the full pipeline against it yields a **generator-conditional** false-positive floor: any signal it produces is definitionally spurious, and memorisation cannot be invoked because there is nothing to memorise. A real candidate’s match must outperform this *one calibrated class of negative controls* by the corrected margin — support means the candidate beat this generator-conditional control family (here a first-order sign-bigram / root-template generator), not that it defeated every possible false-language process. (Its divergence from the calibration targets and the root-template total-variation-distance mechanics are reported in Appendix C.5.)

LLM-ablation compares the pipeline with and without the model’s proposals to isolate contamination — the remaining gate on the contamination number, matching a model’s cognate glosses against the indexed lexical claims.

3.4 The held-out graduation gate

A reading graduates only mechanically, from held-out data, as a conjunction of pre-registered clauses computed by `scripts/verdict.py` — never by the model that proposed the reading. The operative deflation clause is an **order-statistic bar**: the observed held-out statistic must exceed the *expected maximum* of the multiplicity-null taken over the number of trials the search was free to run (the $E[\max]$ over `n_trials` deflation), a **normal-order-statistic approximation** in which $n = n_trials$ is the count of candidates the

search tried, not an estimated number of independent hypotheses. The explicit formula and its Gaussian/independence assumptions are given in Appendix D; the error-control claim is carried not by the formula but by the direct empirical calibration below, and holds conditional on the tested null-generating process, not as a universal guarantee. This is a direct correction for selection across the hypothesis space, and it fails closed on a degenerate null ($\sigma_0 \leq 0 \Rightarrow E[\max] = \mu_0 \Rightarrow$ no graduation). The gate emits one of four mechanical verdicts: **GRADUATE** (every clause holds and the search multiplicity was instrumented), **REJECT** (a clause failed, or the statistic sits below the corrected bar), **NULL_PUBLISHED** (the statistic is indistinguishable from the `L_fake` null — a calibrated non-result, publishable as such), and **NO POWER** (the experiment cannot adjudicate). These verdicts concern *evidential support, not truth*: because Linear A has no ground truth, REJECT / NULL / NO POWER mean a reading is unsupported or insufficiently validated on held-out data, **never that the underlying linguistic proposition is false**. The verdict follows a fixed decision sequence (`verdict.py`): all pre-registered clauses hold on an above-null match $\rightarrow$ **GRADUATE**; else a match clearing the substantive clauses but on an un-instrumented search $\rightarrow$ **INCOMPLETE** (fail-closed); else an observed statistic not exceeding the null mean (indistinguishable from the `L_fake` control) $\rightarrow$ **NULL_PUBLISHED**; else $\rightarrow$ **REJECT**. In the calibration below, the instrumented best-of-100 nulls partition as 3 GRADUATE / 10 NULL_PUBLISHED / 487 REJECT (no INCOMPLETE, the search being instrumented). We calibrate the refusal rule directly: under the tested best-of-100 random-map null — an analyst trying 100 random candidate maps and submitting the best — 3 of 500 replicates graduated, an observed false-graduation rate of 0.6% (one-sided exact Clopper–Pearson 95% upper bound $\approx 1.54\%$), with 487/500 rejected outright and 10 published as calibrated nulls (`scripts/gate_null_calibration.py`); this is an empirical error-control result conditional on that null-generating process, not a universal guarantee. A second pre-registered clause, **generalizes_to_not_indexed**, requires the reading to recover literature-unseen (`L_not_indexed`) signs above a pre-committed

threshold and on a minimum number of distinct such signs — so reproducing already-published values cannot by itself graduate a claim, and the clause fails closed when no threshold or no supported not-indexed sign is present. The candidate must separate from the `L_fake` negative-control distribution, its corrected margin must be cleared, and the observed statistic must beat the order-statistic bar over the Packard/random-lexeme multiplicity null. These are *gate parameters* — pre-registered choices for this corpus, not universal constants. Validation is leave-one-site-out, not ordinary k -fold, because the libation formula induces formulaic dependence across inscriptions. The gate is AND-monotone — clauses can only make it stricter — so hardening never lets a previously-failing reading graduate.

3.5 Grade from persisted artifacts

The invariant across every gate: verdicts are computed from persisted rows, never from a model’s account of its own outcome. The verdict writer is grep-clean of any model call and is the sole writer of the verdicts table. (A self-audit log of confounds this discipline caught inside logos’s own builds is in Appendix D.)

The same falsification-first harness returns *different* calibrated verdicts across probes (Table 1) — the paper’s contribution in one view. Each number is detailed in the cited section.

4 The information floor

Before any decipherment method is graded, the corpus must be sized against what it can identify. Two quantities matter: how much text exists, and how many free parameters a reading spends against it. Confusing the first for the second is the origin of the “too little text” folk claim — and, we argue, of its equally unearned inverse, the optimism that a big-enough corpus makes a reading identifiable.

The corpus. The current authoritative synthesis (Salgarella, 2025) gives an envelope of roughly 2,500 total Linear A finds, of which about 1,400 are readable inscriptions, drawn from 56 find-places spanning ca. 1800–1450 BCE (Middle Minoan II–Late Minoan IB) but overwhelmingly concentrated in the single LM IB destruction horizon (Salgarella, 2025, §1–3, §6). That near-single horizon is why a site-held-out split, not a chronological one, is the only realistic partition. No bilin-

gual text has ever been found (Salgarella, 2025, §8).

Corpus sizing and the unicity computation (numeric detail in Appendix C.4). The recurring sign-count confusion resolves into several distinct sign-count denominators and two occurrence channels that must be kept distinct (enumerated in Appendix C.4). On the narrower parsed syllable-sign budget ($N \approx 5,792$; Salgarella, 2025, §1–3) Shannon’s unicity distance gives $U \approx 204$ –415 signs, everywhere far below N ; we retain this only as a toy-model *precondition* diagnostic and explicitly withdraw it as any sufficiency or refutation claim (the full denominator bookkeeping and its derivation are in Appendix C.4). Qualitatively, the figure cannot settle identifiability either way: unicity distance presupposes a known plaintext-language oracle we do not have, and a lossy, non-injective sign→sound channel — if Linear A underspells like Linear B — may possess no unicity distance at all. Corpus size, in short, cannot decide the question.

What actually binds. The results are consistent with *identifiability* being the more fundamental present constraint rather than corpus size (Salgarella, 2025, §8); we do not establish that more data would be irrelevant. It has three parts: (i) the **absence of a known cognate language** — the map has a value-space but no anchored target (Luo 2019’s limit); (ii) the **multiple-testing / search trap** — that a unique consistent map may exist in principle neither guarantees an algorithm finds it nor stops a cherry-picked false positive across the parameters tried (invariant 8); and (iii) a **mixed inventory** of syllabograms, logograms, and numerals that a real reading must first separate. These, not the count of surviving signs, are what a decipherment must overcome. Accordingly, when a candidate lexicon grows without bound on a fixed occurrence budget — free parameters exceeding the floor — the reading is underdetermined, and the framework must say so.

5 A qualitative audit of the *301 claim

The framework of Section 3 is abstract until it meets a concrete claim. In June 2026 a widely circulated “Linear A cracked” announcement (Di Mino, 2026) argued that Linear A records an extinct Semitic language, reviving Gordon’s hypothesis (Gordon, 1957, 1966). Its load-bearing anchor is a single reading: the recurring peak-sanctuary

Probe	Control / comparator	Real outcome	Null / control floor	Verdict
Graduation gate (§3.4)	—	—	best-of-100 random-map null	3/500 false graduations (0.6%; calibrated under the tested null)
Morphology (§6)	Linear B 9/16	Linear A 4/16	L_fake 6/16 bigram	NO POWER
Segmentation (§6)	—	micro-F1 0.436	random 0.389	supported (gap CI excludes 0)
Metrology (§7.1)	planted 0.9	0.000	permutation 0.113	null (agrees with the field)
Phonology (§7.2)	planted (fires ≥ 2)	0.080 / 0.140	0.065 / 0.201	null (data-limited)
Image alignment (§7.3)	font-swap survives	0.410	image 0.0135 chance	circular demonstration (capped)
Contamination (§8.1)	model-free 0/40	14/40 (JA-NE)	0/40	exploratory (public-exposure gradient)

Table 1: One harness, calibrated epistemic outcomes across probes.

libation formula word A-TA-I-**301*-WA-JA, in which the Linear-A-only sign **301* is *itself* assigned a /na/-initial value, so the verb root reads as West-Semitic N-W-Y “to dwell” (*nawaya*) and is thereafter matched to later Hebrew liturgy. We analyse this **as a claim, on its public evidence** — not the person — from an immutable archived copy carried in the artifact. It is a clean, self-contained instance of the failure modes the gate is built to catch, and the three-part refutation below rests wherever possible on the claim’s **own published table** and on primary sources (Davis, 2013; Duhoux, 1992; Thomas, 2020; Salgarella, 2025; Steele and Meissner, 2017) — none advanced in logos’s own voice, and the paper introduces no new sign reading. We flag the meta-point once: this is a literature-and-primary-source argument, not itself a logos held-out result — the graduation gate and L_fake canary are validated as instruments (unit tests; the known-answer Ugaritic↔Hebrew surrogate) but have **not been run on this claim**. Scoring Gordon’s equations and the Di Mino lexicon against the L_fake floor and the order-statistic bar is directly runnable with the components of §3 and is the priority of the planned revision; that the refutation has not been independently vetted by Aegean epigraphers is a disclosed limitation (§9.6), not a standard the work claims to meet.

(a) The segmentation is not novel. Root *i-**301** plus affixation is the reading of Davis (2013); an independent segmentation by Thomas (2020) reaches the same core and is tested head-to-head against it (the current synthesis, Salgarella, 2025, does not itself adopt Thomas). On the Davis and Thomas segmentations, the affix-stacking is evidenced by multi-attestation forms — ta-na-i-**301*-u-ti-nu (IO Za 6) and a-ta-i-**301*-de-ka (ZA Zb 3) — that vary the material around a stable *i-**301** core. (These sigla are confirmed against Davis 2013: TL Za 1, IO Za 6, and ZA Zb 3.) The claim re-segments a structure already established in the literature. Under the decontamination layer this matters procedurally, not just historically: a correspondence that reproduces a published segmentation is shared-source contamination, tagged `literature_match` and barred from counting as discovery (the artifact, §C.1).

(b) The gloss is contested and weaker. The mainstream value of the verb is “give/dedicate” — Duhoux’s (1992) gloss, adopted by Davis (2013) as a *structural* reading that assigns **301* **no phonetic value at all**, and endorsed by the current synthesis (Salgarella, 2025, §7.2). “Dwell” is a semantically distant, unmotivated substitution. The point is not that “give” is certainly right, but that a *structural* reading needing no sound-value for **301* already accounts for the data the phonetic

claim is invoked to explain.

(c) The primary descriptions of this form’s morphology do not support the Semitic label.

The claim’s own table segments the verb as prefix-stacking and agglutinative — [prefix A “T”] [stem-morpheme TA] [stem-vowel I] [root *301]. We deliberately do **not** rest on an *a-priori* typological contradiction here, and we retract the earlier draft’s claim that a prefix-stacked verb is “internally incoherent” with a Semitic reading: Semitic verbal morphology includes productive prefix conjugations, so a prefix on the verb is not by itself un-Semitic. The point is empirical, not typological, and it comes from the primary literature. Davis (2013, p. 38 n. 13) characterises i-*301’s morphology as “neither IE nor Afroasiatic” — and Semitic *is* Afroasiatic, so a decade earlier the primary description already found **no Afroasiatic morphological signature** in this very form. (We rest on what the footnote states — the absence of an Afroasiatic signature — not on a stronger “ruled out” than a hedged footnote can bear.) Independently, Salgarella (2025, §8) concludes Minoan is agglutinative (after Duhoux 1978) and a language isolate. The claim’s own segmentation is thus consistent with the isolate typology, while the specific Semitic assignment finds no morphological support in the primary descriptions of the form.

The free-parameter receipt. The decisive point for identifiability is that *301 has no cross-script anchor: in the claim’s own table its AB-number column is “—” and its Linear B correspondence “(none)”. Even the standard practice of reading Linear A homomorphs with their Linear B values is a backward projection to be handled with care (Steele and Meissner, 2017, already cautioned by Duhoux 1989, apud Davis 2013, p. 38 n. 11) — and an A-only sign such as *301 has no homomorph to project from at all, so nothing cross-script constrains its value. The single sign on which the entire decipherment pivots is thus a **free parameter** — its value was assigned to fit the desired root, not constrained by any independent evidence. This is precisely the case the minimum-description-length / identifiability *diagnostic* flags: a value chosen to make one word read out carries no description-length saving that survives the multiplicity of alternatives it was selected from (cf. Section 3; the diagnostic is reported next to every claim, and the operative grad-

uation clause — the order-statistic bar over exactly that multiplicity — is one a hand-fitted value cannot clear). For this reading the obstacle is identifiability rather than corpus size — no known cognate language to constrain the assignment, uncorrected multiplicity, and a sign whose value is chosen rather than anchored (a free parameter is unanchored at any corpus size).

Read against Section 3, the claim instantiates two distinct failure modes at once. It is **conclusion-driven model specification** — the table’s *structure* says agglutinative and prefixing, but the *label* says Semitic, and the label wins because the desired conclusion (a Semitic prayer) is a fixed prior rather than an emergent reading. And it is a **free-parameter / identifiability** failure — an unanchored sign-value tuned to the target, on a corpus with no known cognate language to constrain it. Logos indexes the reading (*301 = /na/) precisely so a model cannot regurgitate it as discovery, and quarantines rather than dismisses it. The lesson generalises: a match rate on a hand-picked word, absent a committed prediction and multiplicity accounting, is not evidence. The *301 claim is a qualitative illustration of the failure modes the gate is designed to detect.

6 Results I: morphology induction

We test whether Linear A’s pre-registered affixes and word divisions reflect genuine morphology. The apparent affixation does not clear the operative `L_fake` bigram floor, so the verdict is NO POWER and `has_validated_morphology` = False. The identical test **validates on Mycenaean Greek (Linear B)**, whose longer words carry uncontested inflection: on the 13,562-wordform DĀMOS corpus (mean 3.23 signs/word) it confirms affixation at rate 0.562 — above the same bigram floor (~0.30) — with `has_morphology_power` = True, so the control rules out a universally insensitive detector and supports the reading that Linear A’s shorter forms (mean 1.84) contribute materially to its lack of power (Appendix A). One held-out positive stands and survives its own uncertainty: a DP-unigram segmenter recovers the scribe’s own word divisions at micro-F1 0.436 versus a random-boundary baseline of 0.389, a gap whose site-clustered 95% bootstrap CI [0.021, 0.099] excludes zero ($P(\text{gap} > 0) = 0.998$; the paired-gap detail is in Appendix A). This separates real re-

coverable structure from fitted morphology — the paper’s core discipline.

7 Results II: metrology, phonology, and the cross-script asymmetry

Three further pre-registered probes exercise the harness against distinct field claims. Each returns a calibrated null or a truth-capped positive; together they map where the Linear A corpus does and does not carry recoverable signal, and they are consistent with identifiability rather than effort being the more fundamental present constraint.

7.1 Metrology: a parse-coverage null that agrees with the field

Direction D poses the accounting tablets as a constraint-optimisation: can the fraction-sign values be solved so line items balance the preserved KU-RO totals? Held-out fraction balance does not beat the shuffle null ($p = 1.0$), so the verdict is NULL; the only value the method determines, the fraction sign $J = 1/2$, is the long-standing field consensus (Bennett, 1950), not a logos discovery. This is a null the field itself reached (Corazza et al., 2021) — the harness crediting a surviving value to its originator rather than re-branding it.

7.2 Distributional phonology: a data-limited null bounded by a power curve

Does a sign’s distributional context (which signs sit beside it) carry information about its phonetic value? Tested on the known AB signs, neither the consonant class nor the vowel class clears chance (permutation $p = 0.40 / 0.86$). The verdict is a data-limited NULL: the sound path is not recoverable from Linear A alone, and no sound value is imputed for any sign. It replaces the “just use ML to recover the sounds” objection with a quantified detectability floor.

7.3 The cross-script asymmetry: image alignment recovers shared anchors, but the recovery is circular by construction

Do image embeddings recover the cross-script $A \leftrightarrow B$ correspondence that sequence and cognate co-occurrence cannot? Classical shape descriptors reach direct-NN **0.410** on held-out $A \leftrightarrow B$ anchors while sequence and cognate co-occurrence stay at chance. But the recovery is circular by construction: the $A \leftrightarrow B$ anchor values were themselves assigned by twentieth-century epigraphers on sign-

shape similarity, so this is a shape-genealogy engineering demonstration, truth-layer-capped ≤ 0.75 , not a positive decipherment claim. The study reports no positive decipherment claim.

8 Results III: contamination probes

One failure mode is invisible to the false-positive / multiplicity axis of §7: a decipherment that *looks* novel but silently reproduces published readings a language model has memorised. This section reports the two decontamination probes that target it — both run as method demonstrations, returning completed (if deliberately bounded) findings. Throughout, where a clean number was not obtainable we report the partial-null and the confound rather than a headline.

8.1 LLM-ablation: does a large open-weight model regurgitate the literature?

Does a large open-weight model regurgitate the literature rather than decipher it? In the LLM-ablation, *qwen2.5:72b* reproduces famous published readings — Gordon’s JA-NE = “wine” in 14/40 seeds — that the model-free edit-distance baseline never reaches (0/40), yet never reproduces obscure vessel equations (0/40). Verdict: a reproduction gradient consistent with differential public exposure — though the experiment does not isolate memorisation from correlated item-level differences — the memorisation confound the discipline layer must gate out.

8.2 L_not_indexed: indexed versus not-indexed abstention

Does the model’s behaviour on *published* signs generalise to signs unseen in the indexed literature? *qwen2.5:72b* valued the memorisable known signs in $\sim 20/40$ seeds but abstained $\sim 97\%$ of the time on the literature-unseen *-series signs ($\sim 1/40$) — an abstention asymmetry consistent with differential public exposure (the experiment does not isolate memorisation from correlated item-level differences). With no ground truth on the not-indexed signs, **consistency is not correctness**: it bounds how the model behaves, never whether it is right.

Registered extension: a known-answer sufficiency protocol (asserts no result)

Is a Linear-A-scale corpus large enough to recover a decipherment if a known cognate existed? Be-

cause every benchmark is handed a real cognate signal Linear A lacks, any recovery curve is an **optimistic benchmark, conditional on a known cognate relationship**, readable two ways: at-floor at Linear-A scale $\Rightarrow$ information-insufficient even given a cognate; above-chance $\Rightarrow$ identifiability (cognate-absence, uncorrected multiplicity) rather than corpus size would be implicated. The accuracy-vs-size curve is deposited in the repository as a known-answer upper-bound **protocol**; this paper asserts no interim result from it.

9 Discussion, limitations, and conclusion

9.1 The discipline is the contribution

The deliverable of this work is not a reading of Linear A. It is a *harness* — a literature index with a known-reading versus literature-unseen partition, a fabricated-language control corpus, model-ablation contamination checks, multiple-comparison correction, and pre-registration, all graded mechanically from persisted artifacts — together with a body of pre-registered, calibrated nulls, each with its borderline case named. The stratified morphology re-run is the template: at the pre-registered $\alpha = 0.01$, no pre-registered affix is cross-stratum stable above the bigram floor, robust to the genre-bucketing choice (the H3 nominal near-miss surfacing only in the doubly-relaxed corner is detailed in Appendix A.1). A null that names its threshold, its sensitivity, and its near-miss is a stronger referee object than a flat negative.

The axis this harness targets is the **false-positive / multiplicity axis** — the mechanism by which a determined analyst “translates” almost anything on a small corpus — a property of inference methodology, not of the object or its data (§1–2; see Appendix D for its relation to the survey’s fifteen challenges and to the evaluation problem).

9.2 Two named failure modes for the methods literature

(i) **Narrative capture**. Structure says X; the label says the desired Y; the label wins. In the June-2026 *301 claim the published table segments an agglutinative, prefix-stacking verb, yet the assigned language family is Semitic and the desired output (“prayers to a Goddess in a matriarchal society”) operates as a fixed prior, so the conclusion conforms to the prior rather than emerging from

constrained readings. This is distinct from statistical overfitting and must be named separately. Notably, the current synthesis independently concludes Minoan is agglutinative and a language isolate (Salgarella, 2025, §8), which the segmentation matches while the Semitic label does not — the mismatch is empirical, between the label and the primary descriptions of the form, not an a-priori typological impossibility.

(ii) **Free parameters exceeding the information floor**. A lexicon reported as “408 terms and counting” on a fixed corpus has more degrees of freedom than the data can identify; a decipherment whose lexicon grows weekly is underdetermined by construction. The minimum-description-length / unicity *diagnostic* ($k \leq U_{\text{floor}}$) is built to flag exactly this (it is reported next to every claim rather than enforced as a graduation clause; §3, Appendix D): the single sign the whole reading pivots on carries no cross-script anchor (its AB-number and Linear B columns are empty), so its value was assigned to fit the desired root — the textbook free-parameter case the machinery flags as underdetermined. The refutation rests on the claim’s own published table, archived in the artifact.

9.3 The post-hoc evaluation sequence

These failure modes trace to sequencing. The *301 claim publicly commits to a post-hoc evaluation sequence: declare “officially deciphered,” then vet later. logos’s documented inverse is **pre-register** $\rightarrow$ **held-out** $\rightarrow$ **locked artifact**: a hypothesis is hash-keyed and timestamped before it is tested, graded mechanically against held-out material, and the artifact is fixed before any claim. (The held-out libation-formula corpus and its low-power caveat are detailed in Appendix A.3.)

9.4 Identifiability appears the more fundamental present constraint

We take care not to write a defeatist framing, and equally not to over-claim from corpus size. On the real corpus (1,341 inscriptions; $N \approx 5,792$ parsed syllable-signs, $V = 259$) the toy-model unicity distance is $U \approx 204\text{--}415$ signs — a *precondition* diagnostic only. We explicitly withdraw any reading of this figure as a sufficiency or refutation claim: a lossy, non-injective sign $\rightarrow$ sound channel may have no unicity distance at all, so “ $U \ll N$ ” does **not** license the statement that there is enough text to pin a decipherment (the stance of §4). The

results are consistent with *identifiability* being the more fundamental present constraint — the absence of a known cognate language to map to, an uncorrected multiple-testing search burden, and a mixed sign inventory — rather than a demonstrated corpus-size threshold; a null can arise from identifiability, low power, or size alike, and we do not establish that more data would be irrelevant. Salgarella (2025, §8) reaches the field-level version of this by a different route (§2.2): the etymological method is “fraught with problems” and vocabulary comparison “not probative.”

9.5 Study-wide error budget

Multiplicity discipline operates at two levels, and we are explicit about both. Within a single probe, matches are deflated for the signs $\times$ roots $\times$ hypotheses searched (the attempted-candidate-count correction of §3–4). Across the study, the unit of multiplicity is the **per-probe pre-registration**: the **five executed probes** (morphology, metrology, phonology, cross-script image alignment, and contamination) each carry their own pre-committed α , and every probe is reported regardless of outcome. **Error control is defined within each pre-registered probe; we claim no global family-wise guarantee across probes.** Nothing in the battery is read as a decipherment. (The sufficiency-sweep exclusion, the no-selective-reporting note, and the per-probe reporting-status detail are given in Appendix D.)

9.6 Limitations

The nulls are honestly bounded, not decisive. A dependence-adjusted trial-count correction of the morphology z-test is deferred; the current control is a conservative ≥ 2 -distinct-sites presence gate, which would only strengthen the null (the artifact). No held-out *decipherment* has been achieved, and the study reports **no positive decipherment claim** — the deliverable is calibrated absence, not a reading. Nothing here claims a rejected reading is *false*: with no Linear A ground truth, every REJECT / NULL in this study bounds *evidential support* on held-out data, not truth. The one probe that recovers an above-chance number, cross-script sign-image alignment (direct-NN 0.410 [0.18, 0.64]; Procrustes 0.185 [0.00, 0.36]), is a font-robust but circular *engineering demonstration* (§7.3): a font-swap control refutes the typeface-artifact worry, but the $A \leftrightarrow B$ anchor values were themselves assigned by sign shape, so

the recovery re-derives a shape-defined mapping and is truth-capped at ≤ 0.75 , never a reading. No phonetic claim is made. The philological points in §5 rest on the published literature (Davis, 2013; Duhoux, 1992; Thomas, 2020; Salgarella, 2025; Steele and Meissner, 2017) rather than on independent Aegean-epigraphic review, which we disclose as a limitation; because those points are cited rather than original, the argument does not depend on any new sign reading of our own. Chronological cross-validation is largely foreclosed because roughly 74% of the corpus sits in the single LM IB horizon, leaving site-held-out splits as the only realistic partition. And the corpus itself — small, partly lossy, with no bilingual yet found (Salgarella, 2025, §8) — sets the ceiling.

9.7 Conclusion and reproducibility

The framework is designed to withhold validation when a proposed reading does not clear pre-registered negative controls, multiplicity adjustment, contamination checks, and held-out evaluation; the contribution is the framework and the calibrated null itself — the instrument and the honest limits of what the corpus can support — not a reading, and the calibrated null is a safeguard against unsupported claims. The code, pre-registrations, and all findings are open under the MIT licence and available as an anonymized artifact (code and pre-registrations released on acceptance). The licensed epigraphic corpora — GORILA-, SigLA-, lineara.xyz-, and DĀMOS-derived — are not re-distributed, per their CC BY-NC-SA / all-rights-reserved terms; each source’s licence and retrieval path is documented (data-provenance), so the harness is reproducible against a locally obtained corpus. This work is independent and not affiliated with, or endorsed by, any of the cited projects or their authors.

References

- David H. Bailey and Marcos López de Prado. 2014. The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107.
- E. W. Barber. 1974. *Archaeological Decipherment: A Handbook*. Princeton University Press, Princeton.

- Emmett L. Bennett. 1950. Fractional quantities in Minoan bookkeeping. *American Journal of Archaeology*, 54(3):204–222.
- Taylor Berg-Kirkpatrick and Dan Klein. 2011. Simple effective decipherment via combinatorial optimization. In *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 313–321.
- Taylor Berg-Kirkpatrick and Dan Klein. 2013. Decipherment with a million random restarts. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Jan G. P. Best. 1981. YASSARAM! *Talanta*, 13:17–21.
- Marina Braović, Damir Krstinić, Maja Štula, and Antonia Ivanda. 2024. [A systematic review of computational approaches to deciphering Bronze Age Aegean and Cypriot scripts](#). *Computational Linguistics*, 50(2):725–779.
- Michele Corazza, Silvia Ferrara, Barbara Montecchi, Fabio Tamburini, and Miguel Valério. 2021. [The mathematical values of fraction signs in the Linear A script](#). *Journal of Archaeological Science*, 125:105214.
- Michele Corazza, Silvia Ferrara, Barbara Montecchi, Fabio Tamburini, and Miguel Valério. 2022. [Unsupervised deep learning supports reclassification of the Bronze Age Cypriot writing system](#). *PLoS ONE*, 17(7):e0269544.
- Shruti Daggumati and Peter Z. Revesz. 2019. Data mining ancient scripts to investigate their relationships. In *Proceedings of the 23rd International Database Applications and Engineering Symposium (IDEAS)*.
- Shruti Daggumati and Peter Z. Revesz. 2023. [A method of identifying allographs in undeciphered scripts and its application to the Indus Valley Script](#). *Information*, 14(4):227.
- Herbert A. David and H. N. Nagaraja. 2003. *Order Statistics*, 3rd edition. Wiley, Hoboken, NJ.
- Brent Davis. 2013. [Syntax in Linear A: The word-order of the ‘Libation Formula’](#). *Kadmos*, 52(1):35–52.
- Brent Davis. 2014. *Minoan Stone Vessels with Linear A Inscriptions*. Aegaeum 36. Peeters, Leuven.
- Daniele Di Mino. 2026. “Linear A cracked”: A-TA-I-*301-WA-JA as West-Semitic N-W-Y. Online announcement; no formal venue. Analysed as a public claim from an immutable archived copy (SHA-256 ffc8994224e0288e0d86881a3e98c49b2ffd5aea b5771eb30d30fba0f5608c0e).
- Yves Duhoux. 1978. On the (agglutinative) typology of the Minoan language. Secondary citation, as characterised in Salgarella 2025 §8.
- Yves Duhoux. 1992. Variations morphosyntaxiques dans les textes votifs linéaires A. *Cretan Studies*, 3:65–88.
- Yves Duhoux. 1997. The Linear A locative/allative prefix i-/j- “to/at”. Secondary citation, as cited in Salgarella 2025 §8.
- Nguyen Eu, Duo Xu, and Francesco Perono Cacciafoco. 2019. Coding to decipher Linear A: consonantal comparison against time-compatible lexical lists. In *Proceedings of the 2019 Pacific Neighbourhood Consortium (PNC) Annual Conference and Joint Meetings*, Singapore.
- Silvia Ferrara, Barbara Montecchi, and Miguel Valério. 2021. [What is the ‘Archanes Formula’? deconstructing and reconstructing the earliest attestation of writing in the Aegean](#). *Annual of the British School at Athens*, 116:43–62.
- Silvia Ferrara and Fabio Tamburini. 2022. [Advanced techniques for the decipherment of ancient scripts](#). *Lingue e Linguaggio*, 21(2):239–259.
- Andreas Fuls. 2015. [Classifying undeciphered writing systems](#). *Historische Sprachforschung*, 128(1):42–58.
- I. J. Gelb and R. M. Whiting. 1975. Methods of decipherment. Secondary citation, as characterised in Ferrara and Tamburini (2022, §3.5).
- Louis Godart and Jean-Pierre Olivier. 1976–1985. *Recueil des inscriptions en Linéaire A (GORILA)*. Études crétoises 21. Geuthner, Paris.

- Cyrus H. Gordon. 1957. Notes on Minoan Linear A. *Antiquity*, 31(124):237–240.
- Cyrus H. Gordon. 1966. *Evidence for the Minoan Language*. Ventnor Publishers, Ventnor, NJ.
- Johann-Mattis List. 2014. Investigating the impact of sample size on cognate detection. *Journal of Language Relationship*, 11(1):91–102.
- Christian J. S. Loh and Francesco Perono Cacciafoco. 2020. [A new approach to the decipherment of Linear A, stage 2: A ‘Brute Force Attack’ on an undeciphered writing system](#). In *Grapholinguistics in the 21st Century 2020*, pages 927–943.
- Jiaming Luo, Yuan Cao, and Regina Barzilay. 2019. Neural decipherment via minimum-cost flow: From Ugaritic to Linear B. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 3146–3155.
- Jiaming Luo, Frederik Hartmann, Enrico Santus, Regina Barzilay, and Yuan Cao. 2021. Deciphering undersegmented ancient scripts using phonetic prior. *Transactions of the Association for Computational Linguistics*, 9:69–81.
- David W. Packard. 1974. *Minoan Linear A*. University of California Press, Berkeley and Los Angeles.
- Sujith Ravi and Kevin Knight. 2008. Attacking decipherment problems optimally with low-order n-gram models. In *Proceedings of the 2008 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 812–819.
- Gary A. Rendsburg. 1996. “someone will succeed in deciphering Minoan”: Cyrus H. Gordon and Minoan Linear A. *Biblical Archaeologist*, 59(1):36–43.
- Ester Salgarella. 2019. Drawing lines: The palaeography of Linear A and Linear B. *Kadmos*, 58(1/2):61–92.
- Ester Salgarella. 2020. *Aegean Linear Script(s): Redefining the Relationship between Linear A and Linear B*. Cambridge University Press, Cambridge.
- Ester Salgarella. 2025. *Writing in Bronze Age Crete: “Minoan” Linear A*. Cambridge Elements in Writing in the Ancient World. Cambridge University Press, Cambridge.
- Ester Salgarella and Anna P. Judson. 2024. Signs of the times? testing the chronological significance of Linear A and B palaeography. In *KO-RO-NO-WE-SA (Ariadne Supplement 5)*, pages 359–379.
- Ilse Schoep. 2002. *The administration of neopalatial Crete: a critical assessment of the Linear A tablets*. Minos Supplement 17.
- Peter Schrijver. 2014. [Fractions and food rations in Linear A](#). *Kadmos*, 53(1-2):1–44.
- Benjamin Snyder, Regina Barzilay, and Kevin Knight. 2010. A statistical model for lost language decipherment. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 1048–1057.
- Thea Sommerschildt, Yannis Assael, John Pavlopoulos, Vaughn Stefanak, Andrew Senior, Chris Dyer, John Bodel, Jonathan Prag, Ion Androutsopoulos, and Nando de Freitas. 2023. Machine learning for ancient languages: A survey. *Computational Linguistics*, 49(3):703–747.
- Philippa M. Steele and Torsten Meissner. 2017. From Linear B to Linear A: the problem of the backward projection of sound values. In Philippa M. Steele, editor, *Understanding Relations Between Scripts: The Aegean Writing Systems*, pages 93–110. Oxbow, Oxford.
- Fabio Tamburini. 2025. [On automatic decipherment of lost ancient scripts relying on combinatorial optimisation and coupled simulated annealing](#). *Frontiers in Artificial Intelligence*, 8:1581129. Code: github.com/ftamburin/CSA_OptMatcher.
- R. Thomas. 2020. [Some reflections on morphology in the language of the Linear A libation formula](#). *Kadmos*, 59(1–2):1–23.
- Miguel Valério. 2007. Analysis of Linear A suffixation. Secondary citation, as cited in Salgarella 2025 §8.

Michael Ventris and John Chadwick. 1953. Evidence for Greek dialect in the Mycenaean archives. *Journal of Hellenic Studies*, 73:84–103.

Menghan Zhang and Tao Gong. 2016. [How many is enough? — statistical principles for lexico-statistics](#). *Frontiers in Psychology*, 7:1916.

A Morphology (replication detail)

Replication detail for §6–§8. Nothing in these appendices is load-bearing for the paper’s argument, which stands on the body’s headline nulls; a reader who accepts those may skip them.

A.1 Morphology: rates, floors, word-length, stratification, sensitivity. The probe is pre-registered, unsupervised morphology and segmentation induction over the full Linear A corpus (1,341 inscriptions across 52 sites — the 52 that survive logos normalization of Salgarella’s 56-findplace envelope). Target claims are the recurring affixes and positional grammar of the Davis/Thomas/Duhoux–Valério tradition: prefixes and suffixes (*i-*, *a-*, *-te/-ti*, and Davis’s *i-*301* verb stem). Every candidate was graded against two null floors: a within-word sign-order permutation (shuffle) floor, and an *L_fake* floor built from a first-order sign-bigram (Markov) corpus with zero lexical overlap with the real inscriptions.

Pooled run. Of the 16 pre-registered affixes, 4 beat the within-form permutation null and the deflated multiplicity bar (*i-*301*, *a-*, *i-*, *-te/-ti*), a real confirm rate of 0.250 (Wilson 95% CI [0.10, 0.50]) — clearing the shuffle floor decisively. The apparent standout *i-*301* sits far above the permutation null (six distinct affix environments, $z \approx 8.0$ against the within-form permutation null) — but that figure is computed against the weaker within-form permutation null only, and it does not survive the operative *L_fake* bigram floor. Under the bigram floor the whole set fails: the fabricated Markov corpus confirms the same affixes at rate 0.375 (Wilson 95% CI [0.18, 0.61]), above the real 0.250. The outcome is classified NO POWER, not evidence against morphology; *has_morphology_power* is False and the module reports the null (the Linear-B positive control that validates the detector is reported below).

Word-length distribution. Linear A words are short — mean 1.84 signs over all word-tokens,

median and mode 1, with 76.1% of word-tokens at ≤ 2 signs (56.5% single-sign) — and on such words “morphology” and “bigram order” are statistically indistinguishable. This reconciles with Fuls’s 3.3-sign figure once the denominator is matched: Fuls (2015, 57, §6.5) computes the mean over a word *list* of “554 complete sign sequences, where duplicates are reduced to one” (GORILA; Godart and Olivier, 1976–1985) — i.e. over distinct word *types*, not tokens. Restricting our corpus to distinct multi-sign types likewise recovers 3.07; the token vs. type denominator is the whole discrepancy. Fuls reads the 3.3-sign length, via his exponent-to-word-length regression, as evidence that Minoan is *polysynthetic*; our direct morphology test finds the apparent affixation bigram-reproducible with no power, so the word-length typology and the direct null point opposite ways, and we report the null. The direct positive control — the identical bigram-floor test on Mycenaean Greek (Linear B), whose longer words carry uncontested inflectional morphology — was RUN on the DĀMOS Linear-B corpus (13,562 wordforms, mean 3.23 signs/word, parsed by the same *_damos_wordforms* extractor as §7.3) against a pre-registered 16-affix Mycenaean panel (Ventris and Chadwick, 1953). It fires: REAL confirm rate 0.562 (9/16) — carried by the securely-grammatical suffixes *-jo*, *-o-jo* (gen sg *-oio*), *-de* (allative), *-qe* (enclitic “and”), *-o-i* (dat pl) — against a within-word-shuffle floor ~ 0.03 and an *L_fake* bigram floor ~ 0.30 (stable over five draws), so *has_morphology_power* = True. The identical test that returns NO POWER on Linear A’s short words recovers Mycenaean morphology on longer ones (a sensitivity check: an ad-hoc parser yielding mean 2.19 signs/word did *not* fire — power tracks word length exactly as the short-word argument predicts), which rules out a universally insensitive detector and supports the interpretation that Linear A’s shorter forms contribute materially to its lack of power (*scripts/comparison/linb_morphology_control.py*).

Segmentation positive. Under leave-one-site-out cross-validation (52 sites; not k-fold, given formulaic dependence), a DP-unigram (Goldwater-style) segmenter recovers the scribe’s own word divisions at micro-F1 0.436 versus a random-boundary baseline of 0.389 (boundary base rate 0.406). A site-clustered bootstrap

(resample the 52 sites with replacement, 4,000 draws) puts the real-minus-random gap at a 95% CI of [0.021, 0.099], excluding zero ($P(\text{gap} > 0) = 0.998$); the paired gap, not the wide marginal per-segmenter CIs, is the operative test. Morfessor over-segments the short syllabic corpus (micro-F1 0.156) and is reported as not usable rather than hidden.

Stratified re-run. The strata reconcile exactly to the corpus total: 290 admin + 99 libation + 129 other + 740 seal (structurally excluded as an abbreviation channel) + 83 emptied by abbreviation-strip = 1,341 inscriptions. Induction ran on admin (290 inscriptions, 1,486 words, 14 sites), libation (99, 370, 15), and other (129, 187, 40). The 99-inscription “libation” *genre* stratum — stone-vessel inscriptions over 15 sites — is not the same object as the libation-*formula* corpus (~50 inscribed objects drawn from peak sanctuaries and caves, across more than five findspots). Following Salgarella and Judson (2024), chronological cross-validation was explicitly declined (988/1,341 $\approx$ 74% single-horizon LM IB), and site — not scribal hand, which is not individuated for Linear A (Salgarella, 2019) — is the independence unit, gated by presence at ≥ 2 distinct sites. The result is `has_validated_morphology = False`: no pre-registered affix is cross-stratum stable above the bigram floor at the pre-registered $\alpha = 0.01$. Stratification also showed that even `i-*`301 is bigram-reproducible in the repetitive libation register, where the Markov corpus manufactures the apparent verb morphology; Davis’s slot grammar is not refuted, merely not separable from sign bigrams here. Where formula-bearing sites drive a signal, power is intrinsically low: a leave-one-site-out CV over so few formula findspots (an effective ~5-fold site partition) cannot resolve a modest effect.

Sensitivity, disclosed. The honest 2×2 is: under both α thresholds at the defensible bucketing (libation = stone vessels only), NULL; under loose bucketing (libation including stone objects and inked), NULL at $\alpha = 0.01$ but at $\alpha = 0.05$ the affixes `i-` (locative) and `-te/-ti` (ablative) validate. These are precisely the Duhoux/Valério H3 nominal affixes that Salgarella (2025) §8 independently endorses — the named borderline candidates, fragile to *both* the α threshold and a defensible genre choice. The pre-registered $\alpha = 0.01$ null is itself robust to the bucketing, so the buck-

eting fix is not cherry-picking the null. logos’s own self-review caught the bigram free-rider `a-` and excluded it via the per-affix floor; the adversarial reviewer also corrected a docstring that mislabelled the ≥ 2 -site presence gate as attempted-candidate-count deflation. The present analytical approximation uses attempted-candidate count rather than a dependence-adjusted effective trial count — a recorded, deferred limitation whose correction would only strengthen the null.

A.2 `S_morph` detector calibration (relocated from §3.2). After the productivity-gate fix (§3.2 — an affix must attach to ≥ 2 distinct residual stems), `S_morph`’s measured false-positive rate on within-form-shuffled corpora was **3.0% (6/200; Wilson 95% CI [1.4, 6.4]%)**, consistent with the nominal one-sided $z \geq 2$ rate ($\approx 2.3\%$) within sampling error. We flag this explicitly as an **indicative calibration check, not the operative gate**: a 6/200 estimate is too coarse to certify an $\alpha = 0.01$ threshold on its own, and it is reported only to show the detector is not grossly mis-calibrated at the nominal level. The graduation gate does not rest on it — it applies the order-statistic bar over the far larger permutation and candidate counts the deflation machinery tracks (§3.4). Power was **100% (50/50)** at a single strong planted-affix strength — a saturated point, not a full sensitivity curve (contrast the phonology power curve, §7).

A.3 Held-out design: the libation-formula corpus (relocated from §9.3). The libation-formula corpus (~50 inscribed objects over >5 findspots, spanning peak sanctuaries and caves) offers a natural held-out cross-validation, but a leave-one-site-out split over so few formula sites has low statistical power — and site-level CV is in any case necessary, not sufficient, since within-site scribal-hand correlation must also be controlled (Salgarella and Judson, 2024). This formula corpus should not be conflated with the 99 stone-vessel inscriptions of the “libation” *genre* stratum (§6): they are different objects.

B Metrology, phonology, and cross-script (methods and controls)

B.1 Metrology. Direction D parses HT (LM I) documents into balance units [recipient][commodity logogram][integer][fraction]

signs] and solves the fraction-sign values by RANSAC over exact-Fraction Gaussian elimination (robust to the unbalanced-tablet outliers that defeat least-squares), requiring line items to balance the preserved KU-RO totals; validation is grouped k-fold by document (no tablet straddles train and test) against a permutation null. On the real corpus (32 documents / 35 balance units), integer-level balance holds for 7/28 (25%); of ~ 7 fraction-bearing constraints only one is satisfied by the solved system, and the sole determinable value is $J = 1/2$ (forced by HT104's $2 \cdot J = 1$). Held-out fraction balance is 0.0 against a permutation-null mean of 0.113 (max 0.143), $p = 1.0$ — not separated. A synthetic positive control confirms this is not a machinery bug: with planted values the solver recovers all of them and reaches held-out balance 0.9 vs null 0.22 ($p = 0.003$, separated). The null traces to sparse automated parse coverage (multi-commodity blocks, illegible number glyphs, KI-RO deficit sub-lists, editorial restorations); Corazza et al. (2021, p. 2) abandoned the balance-to-totals method for the same reason. The surviving $J = 1/2$ is the long-standing consensus (Bennett, 1950), with Schrijver's (2014) three further tablets (HT Zd 156, PE 1, HT 123a) lifting it to at least four corroborating documents.

B.2 Distributional phonology. Track 2 asks whether a sign's distributional context — which signs sit beside it — carries information about its phonetic value, measured on the known AB signs whose Linear B sound values can be borrowed. Features are PPMI-weighted, L2-normalised left/right co-occurrence counts; labels are the Linear-B-transferred (C, V) parsed from the GORILA name, so features never see the labels. The test is leave-one-out 1-NN against a label-permutation null, Šidák-deflated over the two tests; 50 of ~ 55 AB signs met the ≥ 3 -occurrence threshold.

Per-test results: consonant class (13 classes) at 0.080 accuracy (null 0.065, permutation $p = 0.40$) and vowel class (5 classes) at 0.140 (null 0.201, permutation $p = 0.86$). The honest verdict is *data-limited*, not *no bridge exists*, because the positive control disciplines the reading. A first draft's flat positive control failed to fire at planted strength 1.0; replacing it with a power curve over planted strengths shows the test firing ($p < .05$) at strength ≥ 2 but not at 1. The test works but

is not infinitely sensitive at $n = 50 / 13$ classes, so we cannot distinguish “no distribution $\rightarrow$ phonetics bridge” from “too few signs to detect a weak one.” Both point the same way: the sound path is not recoverable from Linear A alone. The obvious next move — impute sound values cross-script from a larger Linear B corpus — is now *tested rather than deferred*: §7.3 reports that ingesting the full DĀMOS Mycenaean corpus leaves the cross-script distributional alignment at chance, so more Linear B data is necessary but not sufficient; the block is an absent distributional signal, not corpus volume alone. No sound value is imputed for any sign. This pre-empts the predictable “couldn't you just use ML to recover the sounds?” objection with a quantified detectability floor rather than hand-waving.

B.3 Cross-script image alignment. The palaeographic probe renders 88 Linear-A and 74 Linear-B glyphs and tests held-out recovery of shared $A \leftrightarrow B$ anchors. Classical descriptors (HOG + Hu moments + shape-context) reach direct-NN accuracy 0.410 [0.18, 0.64] and Procrustes 0.185 [0.00, 0.36], far above image chance (0.0135); a from-scratch I-JEPA reaches only 0.324 ± 0.025 (3 seeds) and is excluded from the claims as over-parameterized for ~ 90 glyphs. Image embeddings recover held-out anchors where sequence co-occurrence sat at chance. Two controls decide how to read that asymmetry, and together they *downgrade* it from a positive to a demonstration.

First, a font-swap control. Rendering both scripts in one typeface (Aegean, George Douros) risks making the alignment a single-designer house-style artifact rather than a fact about the signs. We re-render Linear A in Noto Sans Linear A and Linear B in Noto Sans Linear B — two independent type designers, each font covering only its own Unicode block (Noto-A carries zero Linear B glyphs and vice-versa), so no single hand draws both scripts. The recovery survives and is if anything stronger (direct-NN $0.367 \rightarrow 0.416$, Procrustes $0.156 \rightarrow 0.195$ on the control's matched anchor set; both $\gg$ chance): the font-artifact hypothesis is *refuted*, and “distances reflect shape, not font” is now demonstrated rather than asserted (scripts/palaeo/font_control.py).

Second, and decisively, the circularity the font control cannot remove. Linear B was historically adapted from Linear A, and the $A \leftrightarrow B$ transcription *values* that define the anchor set were as-

signed by twentieth-century epigraphers largely on the basis of sign-shape similarity. An image encoder that “recovers the A→B value map” therefore re-derives a correspondence that was *defined* by shape — the result is close to tautological and is not independent evidence for any phonetic reading. We accordingly relabel this leg a **shape-genealogy engineering demonstration, not a positive decipherment claim**: it confirms a known palaeographic fact (borrowed signs look alike) and validates a font-robust descriptor pipeline, but it reads no Minoan. It remains truth-layer-capped in any case — the [Ferrara et al. \(2021\)](#) “Archanes Formula” hard negative (CH 095 resembles AB 60/ra but is not its phonetic counterpart) shows a plausible shape match can be phonetically wrong, so a shape match is a ≤ 0.75 signal (invariant 5), never a reading. The honest asymmetry is thus: the image path re-derives a shape-defined mapping (expected, circular, capped), while the sequence and cognate paths **remain at chance even after ingesting the full DĀMOS Mycenaean corpus** — 13,562 wordforms, $\approx 15\times$ the earlier cognate file, 56 anchors: the best of five alignment methods reaches 0.025 against a 0.011 chance floor (`clearly_above_chance = False`), while the *positive control recovers 0.947*, so the harness works and the null is a real null. Ingesting all of DĀMOS therefore reframes the obstacle from *data volume* to *absent distributional signal*: the A↔B phonetic correspondence is simply not carried by cross-script sign co-occurrence, even at full Linear-B scale. And the A-only signs on which positive-decipherment claims would pivot (e.g. *301) have no cross-script anchor at all, so images cannot impute them even in principle.

C Contamination and sufficiency (detail)

C.1 LLM-ablation reproduction gradient. We ran the §C.4 ablation against `qwen2.5:72b` — a large open-weight model, a far stronger memoriser than a local 12B — on a rented H100 (40 seeds $\times$ 40 held-out Linear A forms; NW-Semitic candidate; a Hebrew lexicon of 7,682 skeletons for the model-free arm; 0 LLM errors across 40/40 seeds) (the artifact). The first-pass contamination rate came out at 1.000 in all 33 seeds where it was defined — but this is an **artifact, reported as such**: the mechanical arm never shares a literature-matching correspondence with

the LLM (their value vocabularies overlap in only 10 of 134 values), so the rate is forced to 1.0 by construction. Publishing it as a contamination result would itself be the overfitting failure the discipline layer exists to prevent (the artifact).

The gate-2 rerun, reconciled to a common consonant-skeleton vocabulary, yields the honest deliverable: a per-word reproduction table over 40 seeds. `qwen2.5:72b` reproduces Gordon’s JA-NE = *yn* / “wine” in 14/40 seeds and Best 1981’s A-SA-SA-RA-ME = “oh Asherah!” in 2/40 — published readings the model-free edit-distance baseline never independently reaches (0/40) — while never reproducing the obscurer vessel equations KA-RO-PA and SU-PA-RA (0/40) (the artifact). Three of the four tested lexical items — JA-NE, KA-RO-PA, SU-PA-RA — are among Gordon’s five West-Semitic equations (as is the separately-treated KU-RO below); the fourth, A-SA-SA-RA-ME, is Best’s, not Gordon’s. An external red-team correction split KU-RO: of its 14/40 hits, 8/40 route through the Semitic root *kull* and 7/40 are reachable via the general-knowledge meaning “total” (confirmed by tablet arithmetic; the two categories overlap on one seed), so KU-RO = “total” is **excluded** as a witness and KU-RO = *kull* is downweighted (the artifact). The surviving gradient — famous readings reproduced, obscure ones not — is what C.4 detects: **a reproduction gradient consistent with differential public exposure**, not isolating memorisation from correlated item-level differences. The gate-2 rate itself (mean 0.646) remains confounded by the baseline’s representational ceiling and is not reported as a contamination measure.

C.2 L_not_indexed abstention asymmetry.

The `L_not_indexed` probe shows a sign in real context and asks whether the model’s behaviour on *published* signs generalises to signs **unseen in the indexed literature** — absent from the indexed sources, hence not memorisable from them (never “impossible to have seen” in any absolute sense) (the artifact). Across 40 seeds, `qwen2.5:72b` assigned a value to the 20 known AB signs in $\sim 20/40$ seeds but to the 11 not-indexed *-series signs in only $\sim 1/40$ — it abstains $\sim 97\%$ of the time on the untouched signs while confidently valuing the memorisable ones. The apparent known–not-indexed consistency gap (0.601, permutation $p = 0.0005$) is coverage-confounded and was flagged, not asserted: the matched- n

robustness check returned `no_power` (0 of 11 not-indexed signs reached $n \geq 10$), so the rigorous comparison could not run. The interpretable signal is the abstention asymmetry itself, consistent with differential public exposure (the experiment does not isolate memorisation from correlated item-level differences) — and, per the module’s standing rule, **consistency is not correctness**: with no ground truth on the not-indexed signs, this bounds how the model behaves, never whether it is right (the artifact). The full arc cost $\approx \$2.80$ in ephemeral GPU rental.

C.3 Information-sufficiency CSA protocol.

The learning-curve harness reframes the withdrawn unicity-distance diagnostic — a toy-model precondition, never a sufficiency or refutation claim — as a sharper, answerable question: *is a Linear-A-scale corpus (539 distinct word-forms; $V = 259$ sign-types at $N \approx 5,792$ parsed syllable-signs) large enough to recover a decipherment IF a known cognate existed?* We report the design and its reading rules here as a known-answer upper-bound protocol; the full accuracy-vs-size curve, with per-size permutation floors, is deposited in and completes in the repository, and this paper asserts no interim result from it.

The protocol downsamples known-answer benchmarks — Linear B / Mycenaean Greek as the primary syllabic analog, sourced through DĀMOS, and further syllabic corpora — and measures held-out cognate-ID recovery against a permutation chance floor, using a published cognate-ID matcher (Tamburini, 2025) that reports the Luo et al. (2019) accuracy metric (the artifact). Because every benchmark is handed a real cognate signal Linear A lacks, any recovery curve is an **optimistic benchmark, conditional on a known cognate relationship** — a best case Linear A does not enjoy — readable two ways: at-floor at Linear-A scale $\Rightarrow$ information-insufficient even given a cognate; above-chance $\Rightarrow$ identifiability — cognate-absence and uncorrected multiplicity over a mixed sign inventory — rather than corpus size would be implicated. Neither branch is ever a decipherability claim, and neither asserts “there is enough text.” Using a data-size recovery curve to bound decipherability is not itself new — decipherment learning curves over cipher length (Ravi and Knight, 2008; Berg-Kirkpatrick and Klein, 2013), the sample-size dependence

of automatic cognate detection (List, 2014), and analytic corpus-sufficiency thresholds (Barber, 1974; Zhang and Gong, 2016) are the ancestors; our contribution is the specific repurposing to a cross-language cognate-ID sufficiency locator for an undeciphered target, not the learning-curve idea.

The matcher is the parallel coupled-annealing path validated in the artifact against its published known-answer benchmarks (a correctness check that the reproduction reaches the published Table-3 values); the accuracy-vs-size curve itself — downsampling each benchmark and reading its recovery against the per-size permutation floor under the two-way logic above — is computed and deposited in the repository, and this paper draws no Linear-A-scale conclusion from it. The phonological / sound path, by contrast, is no longer data-deferred: the DĀMOS Mycenaean holdings are now harvested and ingested, and §7.3 reports that even the full corpus leaves the cross-script distributional alignment at chance — so that bound is a property of the signal, not a promissory note against data yet to arrive. We publish the harness and its bounding logic as an anonymized artifact (released on acceptance).

C.4 Information floor: corpus denominators and the unicity computation (relocated from §4).

The recurring sign-count confusion resolves into four distinct denominators, each of which must be labelled where it is used (Salgarella, 2025, §1–3): **97** (Braović/Tan token-frequent syllabograms), **~150** (full syllabary estimate), **259** (the logos working repertoire, conflating syllabograms, subscripts, and logograms), and **~350** (total including ideograms; GORILA catalogues 390 signs). The channel through which structure is observed is a *sign-occurrence* count, not a document count or a sign inventory (invariant 7), and two occurrence counts must be kept distinct. The broader **~7,400 sign occurrences** across the whole corpus (Schoep 2002; Salgarella, 2025, §1–3) counts every sign type — syllabograms, logograms, and numerals alike. The narrower $N \approx 5,792$ **parsed syllable-signs** is the subset that the unicity computation below actually consumes, after logograms and numerals are stripped. The two are different channel sizes and are used in different places; the **~7,400** figure never enters the unicity arithmetic. Two further properties tighten the budget: Linear A words are short and formulaic (the six-position

libation-formula template; [Salgarella, 2025](#), §7), and a single language/horizon dominates.

On the logos-normalized corpus (1,341 inscriptions, $N \approx 5,792$ parsed syllable-signs, $V = 259$), Shannon’s unicity-distance criterion yields a measured range of $U \approx 204\text{--}415$ **signs** across the whole plausible (V, P) parameter space, everywhere far below N . **We retain this result only as a toy-model precondition diagnostic, and explicitly withdraw it as any sufficiency or refutation claim.** It does *not* show there is enough text to pin a decipherment, and we do not argue that. Unicity distance presupposes a *known plaintext-language oracle*; because we estimate the redundancy D from the ciphertext itself, that oracle collapses into Linear A up to relabelling, so the computation demonstrates only measurable recurrent structure, not that any reading is identifiable. The precondition it checks is necessary, not sufficient. And it can fail to exist entirely: if Linear A underspells like Linear B, the sign→sound channel is lossy and non-injective, and such a channel may possess no unicity distance at all. Corpus size, in short, cannot settle the question either way.

C.5 Literature index and `L_fake` canary (construction detail, relocated from §3.3). *Index seed and adversarial population.* The seed is deliberately conservative — a sign wrongly omitted is wrongly granted `L_not_indexed` status, the dangerous direction — and every disputed entry is indexed to *catch* regurgitation, never as an accepted reading. Population is adversarial: of the seeded West-Semitic lexical proposals, five are Gordon’s equations and a-sa-sa-ra-me = “oh Asherah!” is [Best \(1981\)](#), not Gordon; candidates failing independent corroboration against [Rendsburg \(1996\)](#) (qa-pa = kappu, a Semitic ki-ro) were rejected rather than indexed. Generated from the seed (invariant 12), the index currently holds **63 published claims across 62 indexed entries (55 single signs)** — a deliberately conservative, non-exhaustive seed, not yet the full §C.1 literature census. *`L_fake` calibration.* Its own divergence from the calibration targets is reported, never minimised; with a Semitic root-template sampler and an external Hebrew (ETCBC/bhsa) reject set, the root-template total-variation distance fell from ~ 0.84 to ~ 0.33 end-to-end (~ 0.23 standalone) on the Ugaritic↔Hebrew gold, and any residual real-lexicon collision is measured and published rather than asserted to be zero.

D The order-statistic graduation bar: formula, assumptions, calibration

The operative deflation clause (§3.4) is a normal-order-statistic **approximation** to the expected maximum of the multiplicity-null over the attempted candidates. For a null of mean μ_0 and standard deviation σ_0 and n attempted candidates,

$$E[\max] \approx \mu_0 + \sigma_0 \left[(1 - \gamma) \Phi^{-1} \left(1 - \frac{1}{n} \right) + \gamma \Phi^{-1} \left(1 - \frac{1}{ne} \right) \right]$$

with $\gamma = 0.5772\dots$ the Euler–Mascheroni constant — the standard blended approximation to the mean of the largest of n draws ([David and Nagaraja, 2003](#), *Order Statistics*, 3rd ed.); a candidate clears this clause only if its observed held-out statistic exceeds the bar.

Assumptions and the direction of their error. The approximation is exact only under (i) an approximately Gaussian null and (ii) independence among the n attempted candidates. Neither holds exactly. (i) The permutation / `L_fake` nulls are not perfectly Gaussian; the blended Euler- γ form is a first-order correction, not a guarantee. (ii) The candidates a real search tries are typically **positively correlated** (overlapping sign-value assignments). Positive correlation makes the true expected maximum *smaller* than the independent- n value, so using the raw attempted-candidate count n in place of a smaller effective independent-trial count places the bar *above* the correlated truth — the approximation is therefore **conservative** under positive correlation (it raises the bar and reduces false graduations). We use the raw deduplicated candidate count; a dependence-adjusted effective count is a recorded, deferred refinement (§9.6) whose correction would only strengthen the nulls.

Calibration and its interval. Because the analytic bar rests on assumptions that do not hold exactly, the error-control claim is carried not by the formula but by the direct empirical calibration (§3.4): a Monte-Carlo of the full gate on a best-of-100 random-map null-generating process. The interval on that calibration is the **one-sided exact Clopper–Pearson** bound (the Beta-quantile inversion of the binomial), which needs no normal approximation to the calibration count — for 3 graduations in 500 replicates, a 95% upper bound of $\approx 1.54\%$. [Bailey and López de Prado \(2014\)](#) enter only as selective-inference lineage (§2.4), not as the authority for this bar.

What the gate no longer includes (historical). The minimum-description-length check ($k \leq U_{\text{floor}}$, §4) was, after the same review, **removed from the gate and is now reported as a diagnostic rather than enforced as a graduation clause**: clearing it does not certify identifiability (it is a necessary-precondition / toy-model diagnostic, §4, §9), and as a hard clause it could only be evaluated against a free-parameter default that risked passing degenerate inputs. U_{floor} is still displayed next to every claim, and “parameters > information $\Rightarrow$ underdetermined” is still stated (invariant 7) — it simply no longer sits inside the AND.

Study-wide error budget: elaboration (relocated from §9.5). The sufficiency sweep of the registered sufficiency extension is a known-answer *upper-bound* diagnostic on benchmarks handed a real cognate signal — not a Linear A probe — so it is not among the five executed probes; on the corrected matcher its accuracy-vs-size curve is computed and deposited in the repository (the registered sufficiency extension), and any above-chance recovery on known-answer benchmarks is true by construction, never a Linear A positive requiring multiplicity protection. There is no selective reporting: the one probe that returns an above-chance number — the cross-script image leg (§7.3) — is reported as a circular *engineering demonstration*, not a positive. Accordingly the segmentation result (§6) is reported as a *probe-level* positive, and the cross-script image leg (§7.3, capped at ≤ 0.75 and relabeled circular) as a capped demonstration — neither is asserted as a battery-level-protected positive.

Baseline and self-audit provenance (relocated from §3.4–3.5). The match-rate the deflation is applied to is supplied by a cognate-ID baseline that reports the Luo et al. (2019) accuracy metric; the instance we run is Tamburini’s CSA_OptMatcher (coupled simulated annealing, §2.3–2.4, reproduced in the artifact), so the gate does not depend on any single citation. The persisted-artifact discipline caught confounds inside logos’s own builds — a S_{morph} statistic that conflated repetition with morphology, a verdict clause that was a tautology and could never fire False, and a search-log tripwire that flagged honest multi-sign logs. Each was fixed and locked by regression tests before commit, illustrating the harness turned on itself.

The axis, relative to the survey (relocated from §9.1). As set out in the Introduction, the false-positive / multiplicity axis is a property of inference methodology and epistemics, not of the object or its data; it is complementary to, not a member of, the fifteen data-and-object obstacles enumerated by Braović et al. (2024), and distinct from the evaluation problem their §6.3 gestures at.

Software & data. The artifact. Code, pre-registrations, and findings are released as an anonymized artifact on acceptance. Licensed corpora (GORILA-, SigLA-, lineara.xyz-, DĀMOS-derived) are not redistributed; a data-provenance note documents each source’s licence and retrieval path.